

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.ticker as tcr
%matplotlib inline
import seaborn as sns
```

In [2]:

```
#reading data as excel as csv file is too big
dating = pd.read_excel('/home/amybirdee/hobby_projects/dating_site/profiles.xlsx')
```


In [3]:

```
dating.head(3)
```

Out[3]:

	age	body_type	diet	drinks	drugs	education	essay0	
0	22	a little extra	strictly anything	socially	never	working on college/university	removed_for_privacy	remove
1	35	average	mostly other	often	sometimes	working on space camp	removed_for_privacy	remove
2	38	thin	anything	socially	NaN	graduated from masters program	removed_for_privacy	remove

3 rows × 31 columns



In [4]:

```
#checking number of rows and columns
dating.shape
```

Out[4]:

(59946, 31)

In [5]:

```
#checking datatypes
dating.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59946 entries, 0 to 59945
Data columns (total 31 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   59946 non-null  int64
 1   body_type             54650 non-null  object
 2   diet                  35551 non-null  object
 3   drinks               56961 non-null  object
 4   drugs                45866 non-null  object
 5   education             53318 non-null  object
 6   essay0               59946 non-null  object
 7   essay1               59946 non-null  object
 8   essay2               59946 non-null  object
 9   essay3               59946 non-null  object
10  essay4               59946 non-null  object
11  essay5               59946 non-null  object
12  essay6               59946 non-null  object
13  essay7               59946 non-null  object
14  essay8               59946 non-null  object
15  essay9               59946 non-null  object
16  ethnicity             54263 non-null  object
17  height                59940 non-null  float64
18  income                59943 non-null  float64
19  job                   51745 non-null  object
20  last_online           59943 non-null  object
21  location              59943 non-null  object
22  offspring             24383 non-null  object
23  orientation           59943 non-null  object
24  pets                  40023 non-null  object
25  religion              39717 non-null  object
26  sex                   59943 non-null  object
27  sign                  48887 non-null  object
28  smokes                54431 non-null  object
29  speaks                59893 non-null  object
30  status                59943 non-null  object
dtypes: float64(2), int64(1), object(28)
memory usage: 14.2+ MB
```

In [6]:

```
#describing the data
dating.describe(include = 'all')
```

Out[6]:

	age	body_type	diet	drinks	drugs	education	essay
count	59946.000000	54650	35551	56961	45866	53318	59946.000000
unique	NaN	12	18	6	3	32	NaN
top	NaN	average	mostly anything	socially	never	graduated from college/university	removed_for_privacy
freq	NaN	14652	16585	41780	37724	23959	59946.000000
mean	32.340290	NaN	NaN	NaN	NaN	NaN	NaN
std	9.452779	NaN	NaN	NaN	NaN	NaN	NaN
min	18.000000	NaN	NaN	NaN	NaN	NaN	NaN
25%	26.000000	NaN	NaN	NaN	NaN	NaN	NaN
50%	30.000000	NaN	NaN	NaN	NaN	NaN	NaN
75%	37.000000	NaN	NaN	NaN	NaN	NaN	NaN
max	110.000000	NaN	NaN	NaN	NaN	NaN	NaN

11 rows × 31 columns

In [7]:

```
#checking column names
dating.columns
```

Out[7]:

```
Index(['age', 'body_type', 'diet', 'drinks', 'drugs', 'education', 'essay0',
      'essay1', 'essay2', 'essay3', 'essay4', 'essay5', 'essay6', 'essay7',
      'essay8', 'essay9', 'ethnicity', 'height', 'income', 'job',
      'last_online', 'location', 'offspring', 'orientation', 'pets',
      'religion', 'sex', 'sign', 'smokes', 'speaks', 'status'],
      dtype='object')
```

In [8]:


```
#dropping the essay columns and last_online as not needed for analysis
dating = dating.drop(['essay0', 'essay1', 'essay2', 'essay3', 'essay4', 'essay5', 'essay6', 'essay7', 'essay8', 'essay9', \
                    'last_online'], axis = 1)
```

In [9]:

```
dating.head()
```

Out[9]:

	age	body_type	diet	drinks	drugs	education	ethnicity	height	income
0	22	a little extra	strictly anything	socially	never	working on college/university	asian, white	75.0	-1.0
1	35	average	mostly other	often	sometimes	working on space camp	white	70.0	80000.0
2	38	thin	anything	socially	NaN	graduated from masters program	NaN	68.0	-1.0
3	23	thin	vegetarian	socially	NaN	working on college/university	white	71.0	20000.0
4	29	athletic	NaN	socially	never	graduated from college/university	asian, black, other	66.0	-1.0



In [10]:

```
#filling in nan values - using median for age since dataset is right skewed and average  
s for rest of numeric values.  
#Using most common values for some non-numeric variables where there are only a few mis  
sing values.  
#Otherwise filling with 'no response given'  
dating['age'] = dating.age.fillna(dating['age'].median())  
dating['body_type'] = dating.body_type.fillna('average')  
dating['diet'] = dating.diet.fillna('no response given')  
dating['drinks'] = dating.drinks.fillna('no response given')  
dating['drugs'] = dating.drugs.fillna('no response given')  
dating['education'] = dating.education.fillna('no response given')  
dating['ethnicity'] = dating.ethnicity.fillna('no response given')  
dating['height'] = dating.height.fillna(dating['height'].mean())  
dating['income'] = dating.income.fillna(dating['income'].mean())  
dating['job'] = dating.job.fillna('no response given')  
dating['location'] = dating.location.fillna('no response given')  
dating['offspring'] = dating.offspring.fillna('no response given')  
dating['orientation'] = dating.orientation.fillna('no response given')  
dating['pets'] = dating.pets.fillna('no response given')  
dating['religion'] = dating.religion.fillna('no response given')  
dating['sex'] = dating.sex.fillna('no response given')  
dating['sign'] = dating.sign.fillna('no response given')  
dating['smokes'] = dating.smokes.fillna('no response given')  
dating['speaks'] = dating.speaks.fillna('no response given')  
dating['status'] = dating.status.fillna('no response given')
```

In [11]:

```
#apostrophes are being replaces with '&rsquo;' - fixing this  
dating['offspring'] = dating['offspring'].str.replace('doesn&rsquo;t', "doesn't")  
dating['sign'] = dating['sign'].str.replace('doesn&rsquo;t', "doesn't")
```

In [12]:

```
#converting numeric fields to integers rather than floats  
dating['age'] = dating.age.astype(int)  
dating['height'] = dating.height.astype(int)  
dating['income'] = dating.income.astype(int)
```

In [13]:

```
#all records now filled in  
dating.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 59946 entries, 0 to 59945  
Data columns (total 20 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   age             59946 non-null  int64  
1   body_type       59946 non-null  object  
2   diet            59946 non-null  object  
3   drinks         59946 non-null  object  
4   drugs          59946 non-null  object  
5   education       59946 non-null  object  
6   ethnicity       59946 non-null  object  
7   height         59946 non-null  int64  
8   income         59946 non-null  int64  
9   job            59946 non-null  object  
10  location        59946 non-null  object  
11  offspring       59946 non-null  object  
12  orientation     59946 non-null  object  
13  pets           59946 non-null  object  
14  religion        59946 non-null  object  
15  sex            59946 non-null  object  
16  sign           59946 non-null  object  
17  smokes         59946 non-null  object  
18  speaks         59946 non-null  object  
19  status         59946 non-null  object  
dtypes: int64(3), object(17)  
memory usage: 9.1+ MB
```

In [14]:

```
#checking age distribution - shows a right skew
plt.figure(figsize = (8,8))
ax = plt.subplot()

dating['age'].hist(bins = 40, color = 'red')

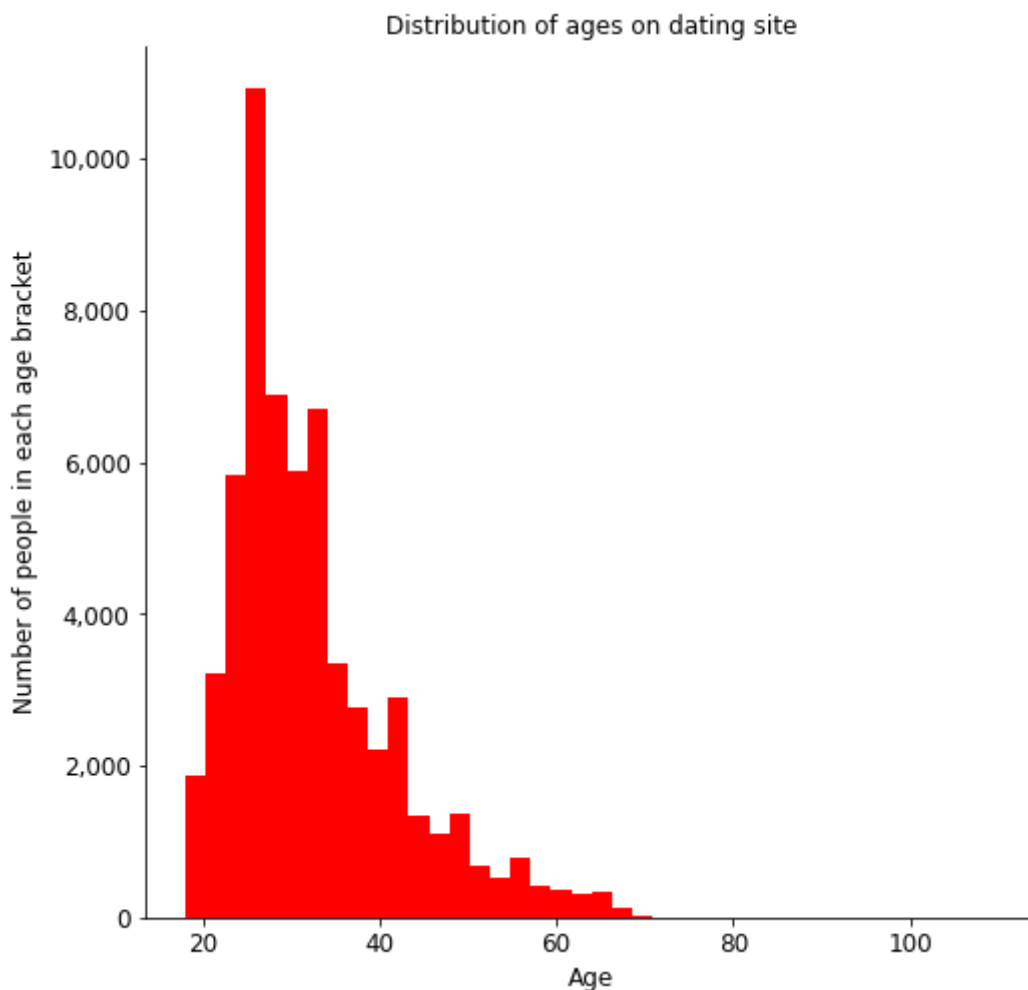
#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#function to add comma separator to labels. Function takes tick label and tick position
def comma(x, pos):
    return format(x, ",.0f")

#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Age', fontsize = 12)
plt.ylabel('Number of people in each age bracket', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Distribution of ages on dating site', fontsize = 12)
plt.grid(None)
plt.savefig('Age - histogram', bbox_inches = 'tight')
```



In [15]:

```
#grouping by age to see the data - a couple of ages over 100 - anomalies or might be false data  
Age = dating.groupby('age').size().reset_index().rename(columns = {0: 'count_of_age'})  
Age.head()
```

Out[15]:

	age	count_of_age
0	18	309
1	19	611
2	20	953
3	21	1282
4	22	1934

In [16]:

```
#plotting a scatter plot of ages - most people are mid 20s-30
plt.figure(figsize = (6,6))
ax = plt.subplot()


plt.scatter(Age['age'], Age['count_of_age'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

#this code adds a comma separator to the y tick marks
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Age', fontsize = 12)
plt.ylabel('Number of people', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Count of age for dating site members', fontsize = 12)
plt.savefig('Age - scatter plot')
```



In [17]:

```
#grouping ages into 7 groups for barplot using pd.cut to cut the age column
bins = [17, 19, 29, 39, 49, 59, 69, np.inf]
labels = ['<20', '20-29', '30-39', '40-49', '50-59', '60-69', '70+']
Age['age_range'] = pd.cut(Age['age'], bins = bins, labels = labels)
Age.head()
```

Out[17]:

	age	count_of_age	age_range
0	18	309	<20
1	19	611	<20
2	20	953	20-29
3	21	1282	20-29
4	22	1934	20-29

In [18]:

```
#grouping age ranges
age_range = Age.groupby('age_range').count_of_age.sum().to_frame().reset_index()
age_range
```

Out[18]:

	age_range	count_of_age
0	<20	920
1	20-29	27821
2	30-39	19846
3	40-49	7338
4	50-59	2860
5	60-69	1159
6	70+	2

In [19]:


```
#barplot of ages - ci = None removes the confidence intervals
plt.figure(figsize = (8,8))
ax = sns.barplot(x = age_range['age_range'], y = age_range['count_of_age'], color = 'grey', ci = None)

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Age range', fontsize = 12)
plt.ylabel('Count of age range', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Dating site members by age range', fontsize = 12)
plt.savefig('Age - barplot', bbox_inches = 'tight')
```



In [20]:

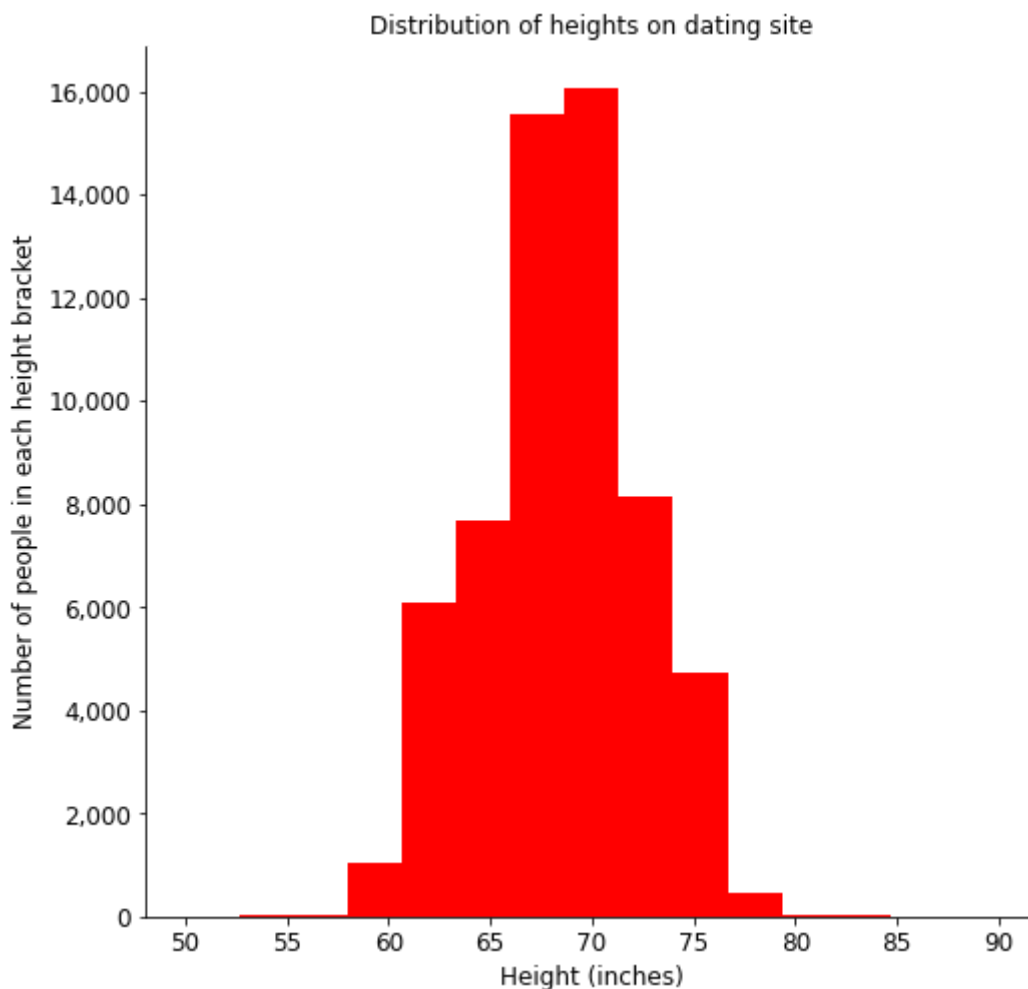
```
#checking height distribution - shows a relatively normal distribution
plt.figure(figsize = (8,8))
ax = plt.subplot()
dating['height'].hist(bins = 15, range = [50, 90], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Height (inches)', fontsize = 12)
plt.ylabel('Number of people in each height bracket', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Distribution of heights on dating site', fontsize = 12)
plt.grid(None)
plt.savefig('Height - histogram', bbox_inches = 'tight')
```



In [21]:

```
#grouping by height to see the data - most people are 60-70 inches but 20 are 95 inches. A few people in the 1-9 inch  
#category - these are likely incorrect datapoints which haven't been filled in properly  
height = dating.groupby('height').size().reset_index().rename(columns = {0: 'count_of_height'})  
height.head()
```

Out[21]:

	height	count_of_height
0	1	1
1	3	1
2	4	1
3	6	1
4	8	1

In [22]:

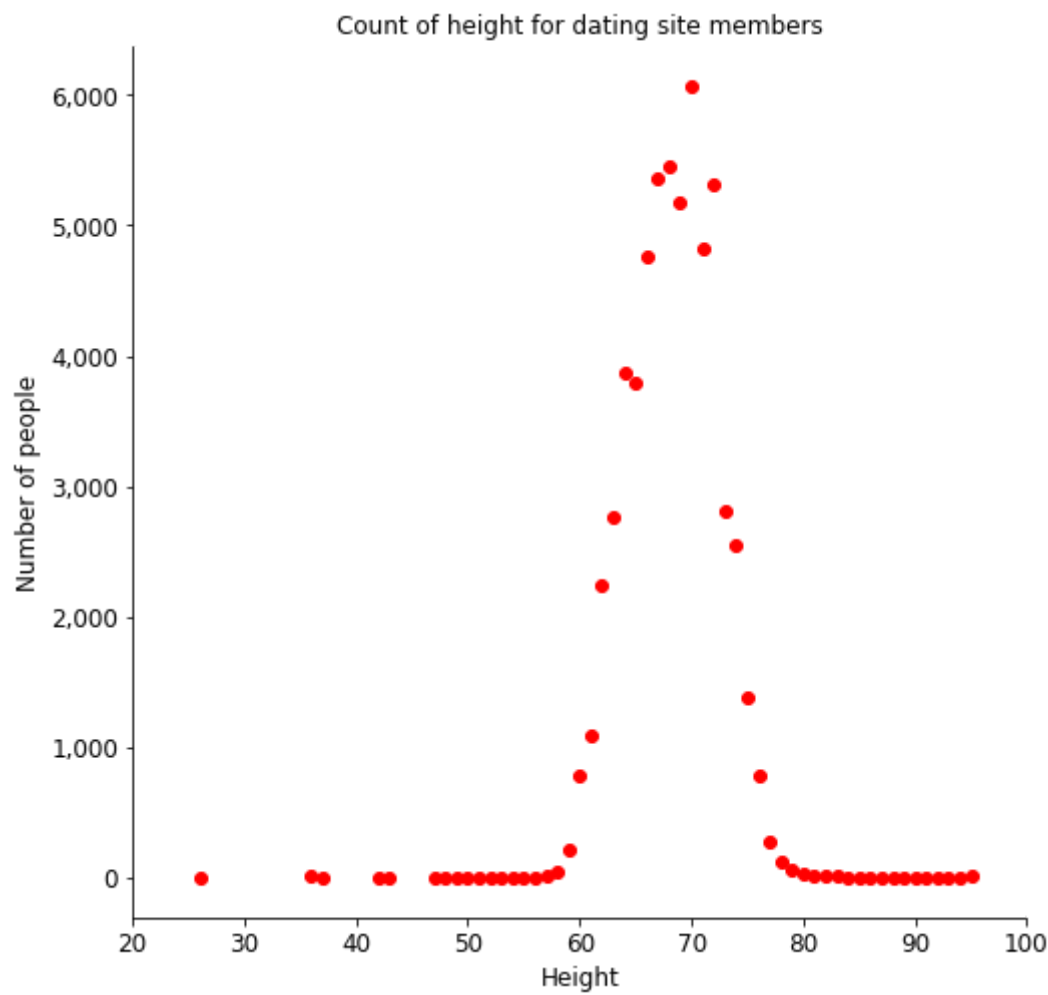
```
#plotting a scatter plot of heights - have excluded the very low heights that were clearly wrong
plt.figure(figsize = (8,8))
ax = plt.subplot()
plt.scatter(height['height'], height['count_of_height'], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

ax.set_xlim(20, 100)
plt.xlabel('Height', fontsize = 12)
plt.ylabel('Number of people', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Count of height for dating site members', fontsize = 12)
plt.savefig('Height - scatter plot')
```



In [23]:

```
#grouping heights into groups for barplot using pd.cut to cut the height column  
bins = [0, 40, 50, 60, 70, 80, 90, np.inf]  
labels = ['<40', '40-50', '50-60', '60-70', '70-80', '80-90', '90+']  
height['height_range'] = pd.cut(height['height'], bins = bins, labels = labels)  
height.head()
```

Out[23]:

	height	count_of_height	height_range
0	1	1	<40
1	3	1	<40
2	4	1	<40
3	6	1	<40
4	8	1	<40


In [24]:

```
#barplot of heights
plt.figure(figsize = (8,8))
ax = sns.barplot(x = height['height_range'], y = height['count_of_height'], color = 'grey', ci = None)

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Height range', fontsize = 12)
plt.ylabel('Count of height range', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)
plt.title('Dating site members by height range', fontsize = 12)
plt.savefig('Height - barplot')
```



In [25]:

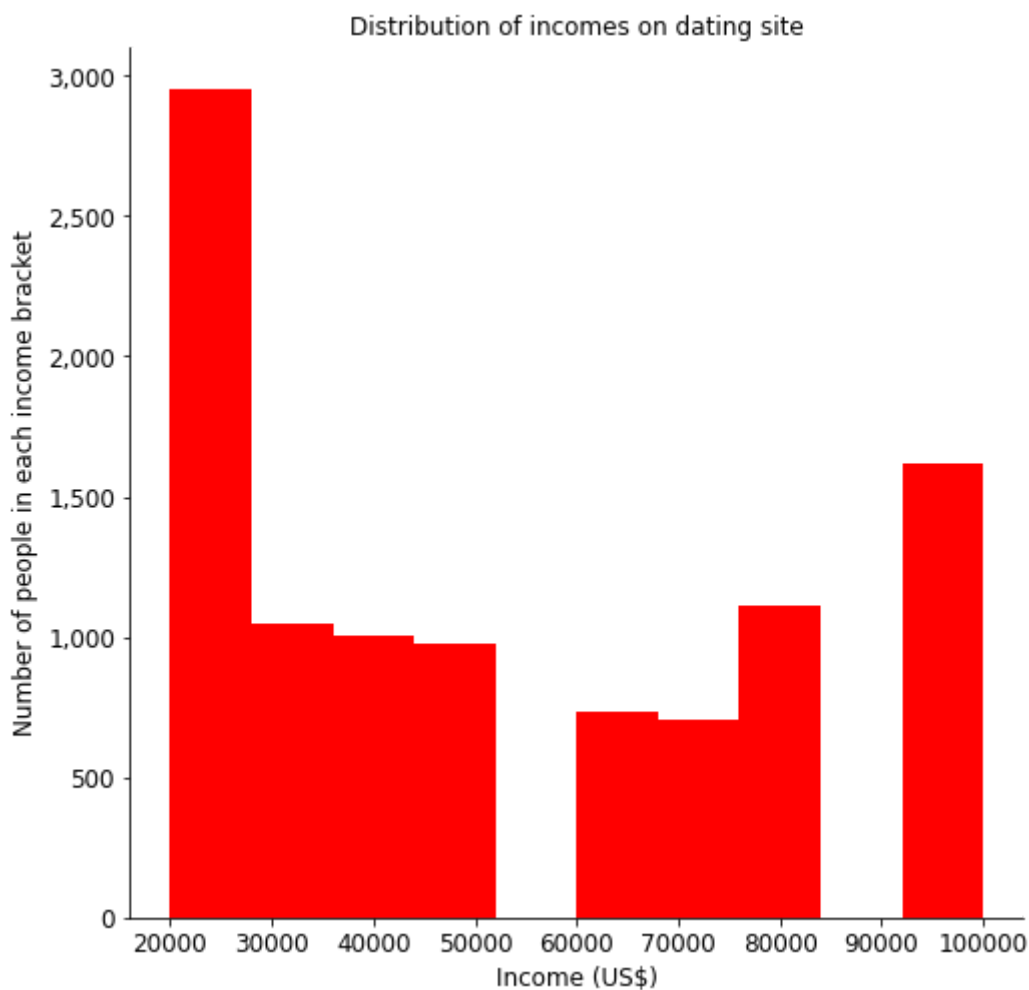
```
#checking income distribution - most people are on lower incomes but many did not fill
this fields in. A few people also
#said their income was $1,000,000 which propbably isn't true - have excluded from chart
plt.figure(figsize = (8,8))
ax = plt.subplot()
dating['income'].hist(bins = 10, range = [20000, 100000], color = 'red')

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Income (US$)', fontsize = 12)
plt.ylabel('Number of people in each income bracket', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Distribution of incomes on dating site', fontsize = 12)
plt.grid(None)
plt.savefig('Income - histogram')
```



In [26]:

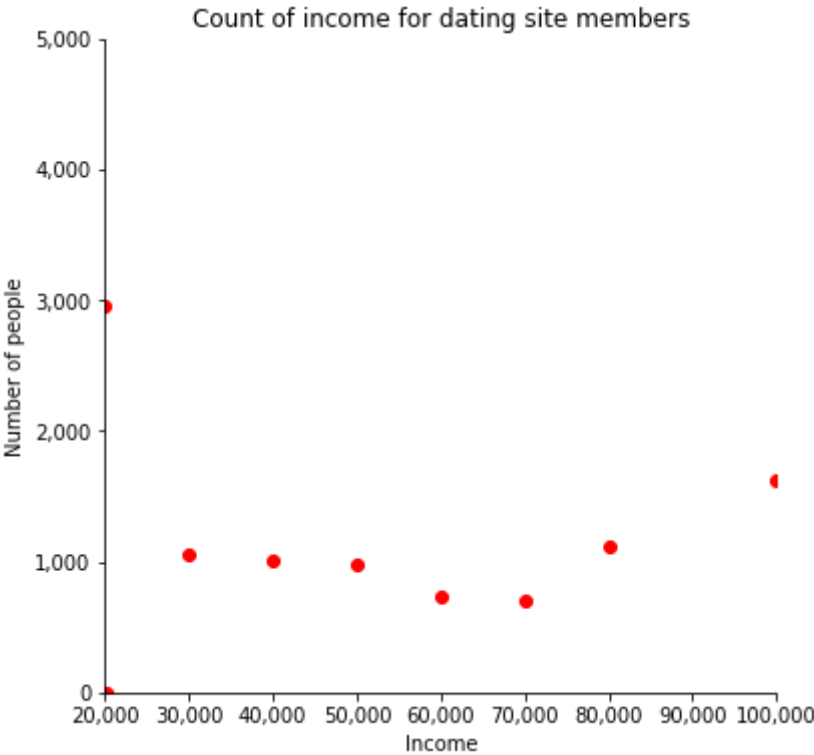
```
#grouping by income to see the data - 81% of people put -1 for this which likely means they didn't answer the question  
income = dating.groupby('income').size().reset_index().rename(columns = {0: 'count_of_income'})  
income.head()
```

Out[26]:

	income	count_of_income
0	-1	48440
1	20000	2952
2	20033	3
3	30000	1048
4	40000	1004

In [27]:

```
#plotting a scatter plot of income - have excluded the very low and high incomes that w  
ere clearly wrong or incomplete  
plt.figure(figsize = (6,6))  
ax = plt.subplot()  
plt.scatter(income['income'], income['count_of_income'], color = 'red')  
  
#removing chart borders  
ax.spines['top'].set_visible(False)  
ax.spines['right'].set_visible(False)  
  
ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))  
ax.xaxis.set_major_formatter(tcr.FuncFormatter(comma))  
  
ax.set_xlim(20000, 100000)  
ax.set_ylim(0, 5000)  
plt.xlabel('Income')  
plt.ylabel('Number of people')  
  
plt.title('Count of income for dating site members')  
plt.savefig('Income - scatter plot')
```



In [28]:

```
#grouping income into groups for barplot using pd.cut to cut the income column  
bins = [0, 40000, 60000, 80000, 100000, np.inf]  
labels = ['<40', '40-60', '60-80', '80-100', '100+']  
income['income_range'] = pd.cut(income['income'], bins = bins, labels = labels)  
income
```

Out[28]:

	income	count_of_income	income_range
0	-1	48440	NaN
1	20000	2952	<40
2	20033	3	<40
3	30000	1048	<40
4	40000	1004	<40
5	50000	975	40-60
6	60000	736	40-60
7	70000	707	60-80
8	80000	1111	60-80
9	100000	1621	80-100
10	150000	631	100+
11	250000	149	100+
12	500000	48	100+
13	1000000	521	100+

In [29]:

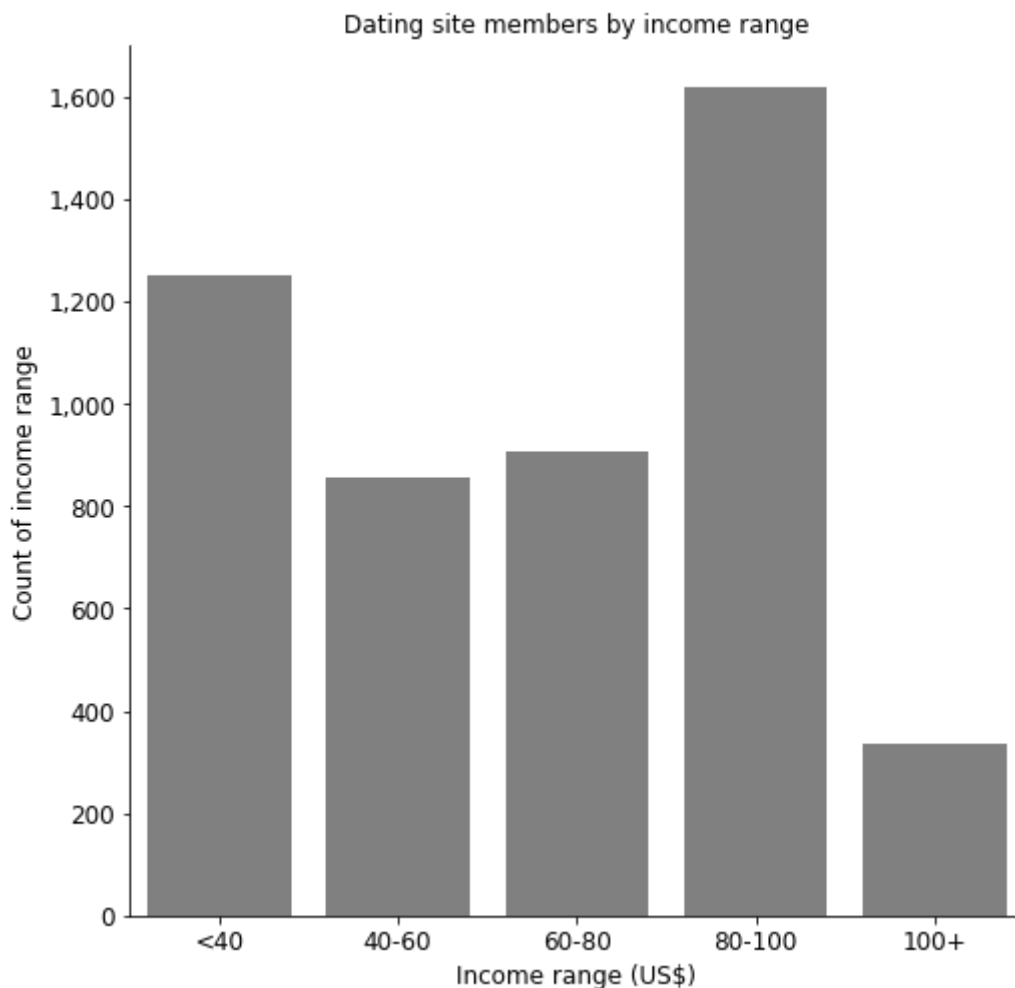
```
#barplot of income - have excluded those who didn't fill this in
plt.figure(figsize = (8,8))
ax = sns.barplot(x = income['income_range'], y = income['count_of_income'], color = 'grey', ci = None)

#removing chart borders
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

ax.yaxis.set_major_formatter(tcr.FuncFormatter(comma))

plt.xlabel('Income range (US$)', fontsize = 12)
plt.ylabel('Count of income range', fontsize = 12)
plt.tick_params(axis = 'x', labelsize = 12)
plt.tick_params(axis = 'y', labelsize = 12)

plt.title('Dating site members by income range', fontsize = 12)
plt.savefig('Income - barplot')
```



In [30]:

```
#converting dating dataframe to csv as will continue analysing categorical values in  
a new notebook  
dating = dating.to_csv('/home/amybirdee/hobby_projects/dating_site/dating_clean.csv', i  
ndex = False)
```

In []: